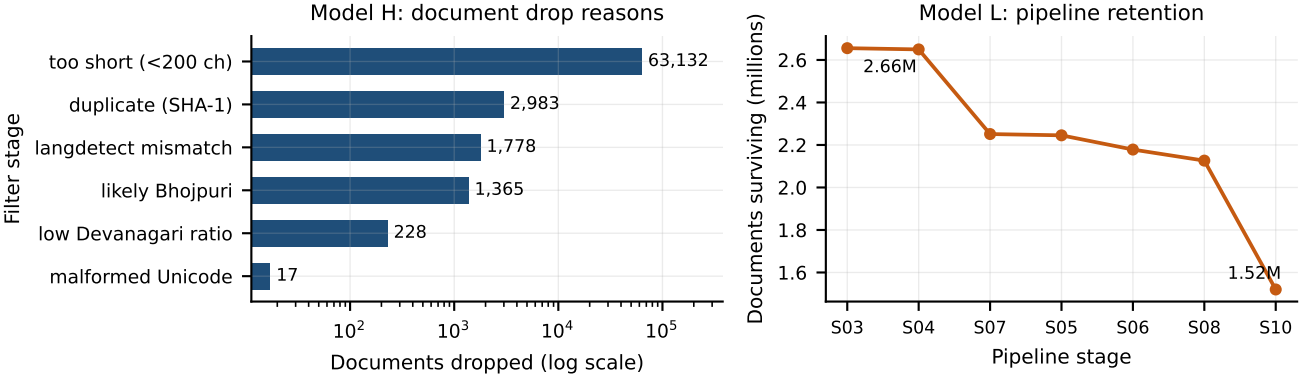
Model H: document drop reasons
too short (<200 ch)
duplicate (SHA-1)
langdetect mismatch
likely Bhojpuri
low Devanagari ratio
malformed Unicode
63,132
2,983
1,778
1,365
228
17
Filter stage
10^2
10^3
10^4
10^5
Documents dropped (log scale)
Model L: pipeline retention
Documents surviving (millions)
2.6
2.4
2.2
2.0
1.8
1.6
2.66M
1.52M
S03
S04
S07
S05
S06
S08
S10
Pipeline stage